



Enable and Configure Disaster Recovery (DR) Replication




What is Vault Replication?

Organizations usually have infrastructure that **spans multiple datacenters**

- Vault needs to be highly-available for application access
- Needs to scale as organizations continue to add use cases and apps
- Common set of policies that are enforced **globally**
- Consistent set of secrets and configurations available to applications that need them regardless of data center



What is Vault Replication?



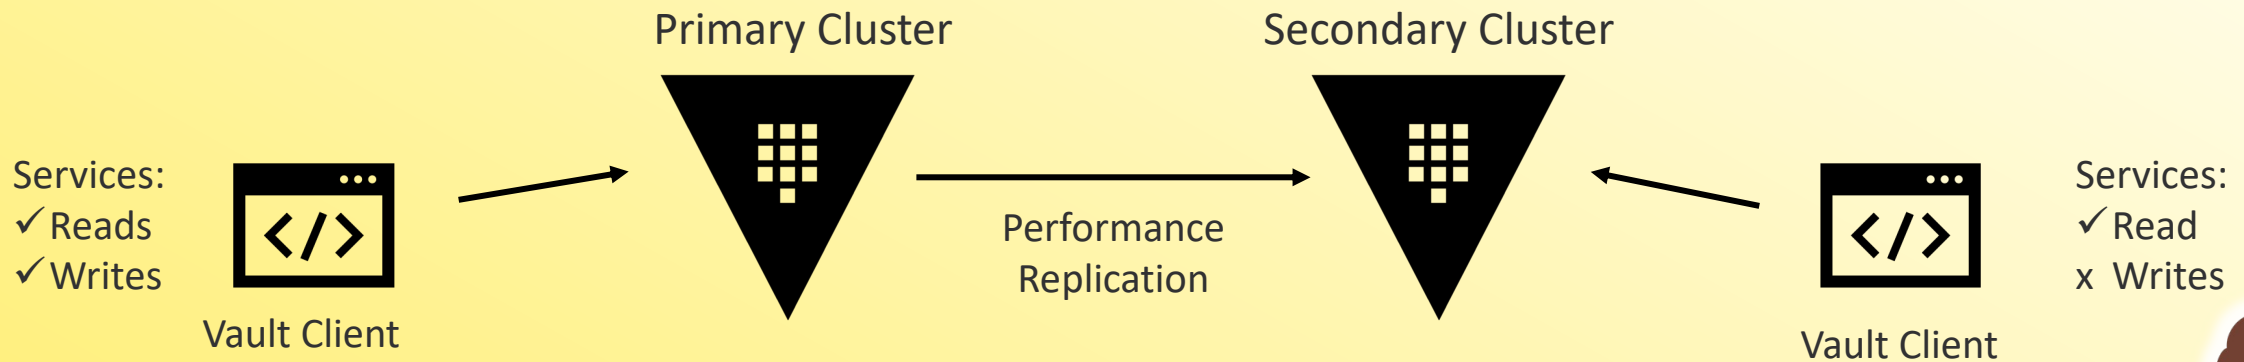
- Only available in Vault Enterprise
- Replication operates on a leader-follower model (**primaries** and **secondaries**)
- The primary cluster acts as the system of record and replicates most Vault data asynchronously
- All communication between primaries and secondaries is **end-to-end encrypted** with mutually-authenticated TLS sessions



Performance Replication



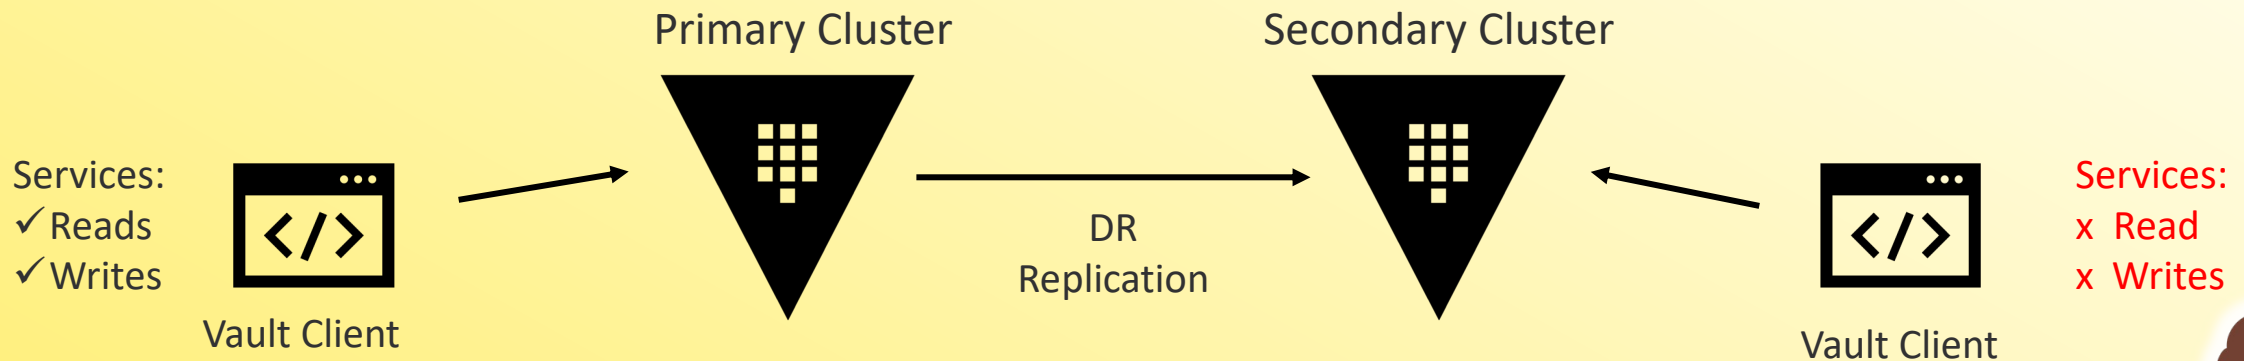
- Replicates the underlying configuration, policies, and other data
- Ability to service reads from client requests
- Clients will authenticate to the performance replicated cluster separately
- Does not replicate tokens or leases to performance secondaries



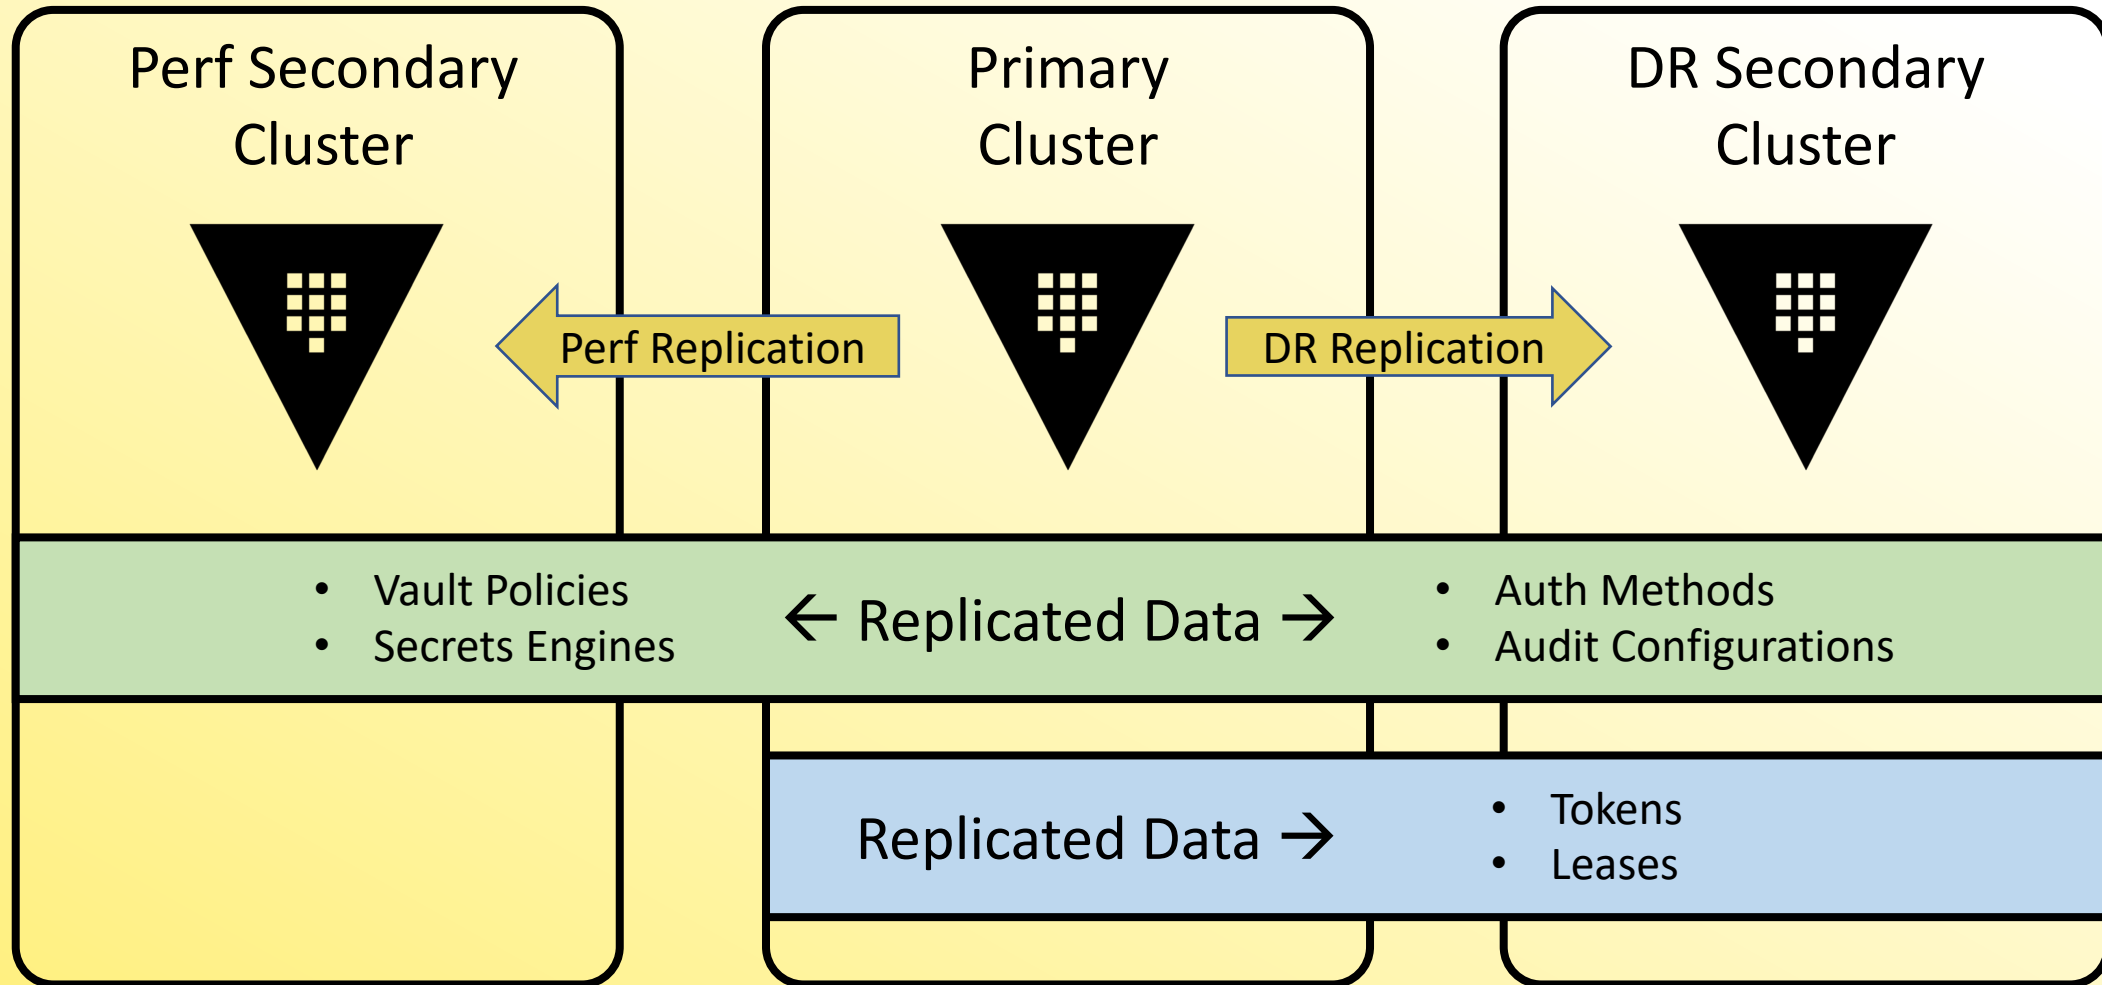
Disaster Recovery Replication



- Replicates the underlying configuration, policies, and all other data
- Cannot service reads from client requests
- Clients should authenticate with the primary cluster only (or a perf cluster)
- Will replicate tokens and leases created on the primary cluster



Comparison



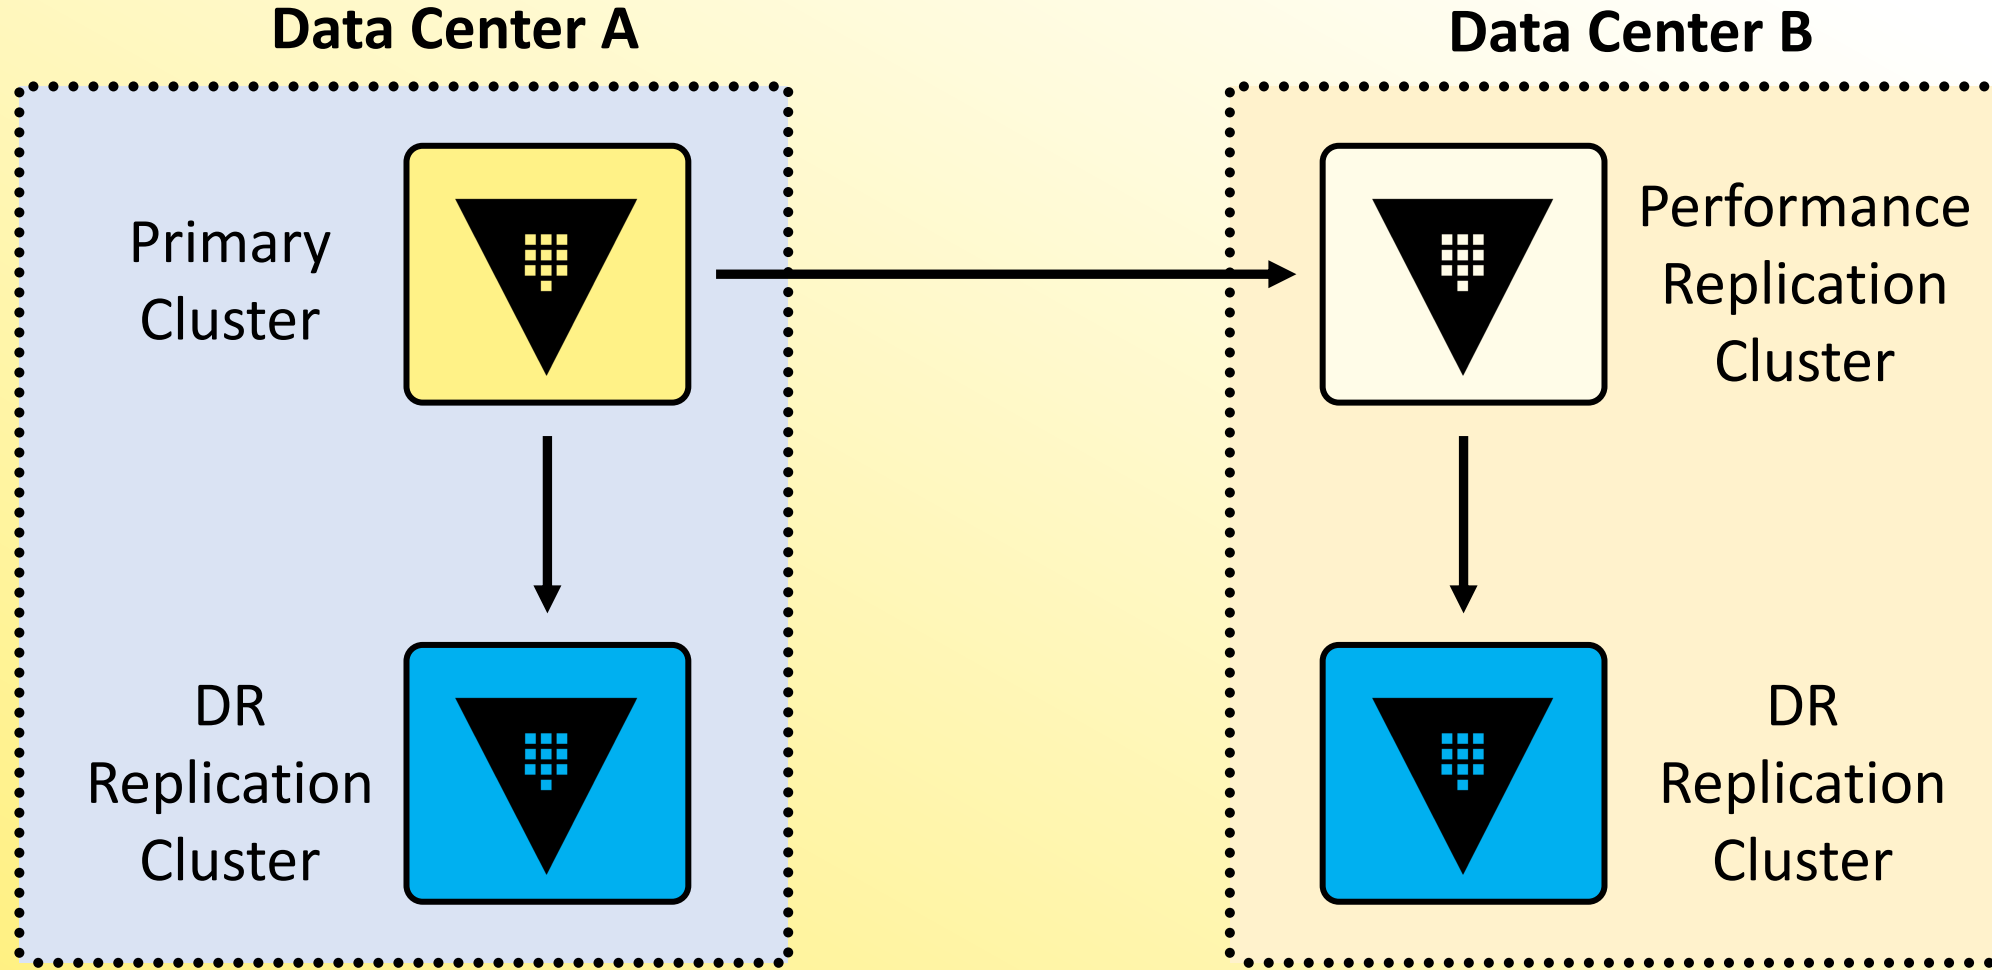
Disaster Recovery Replication



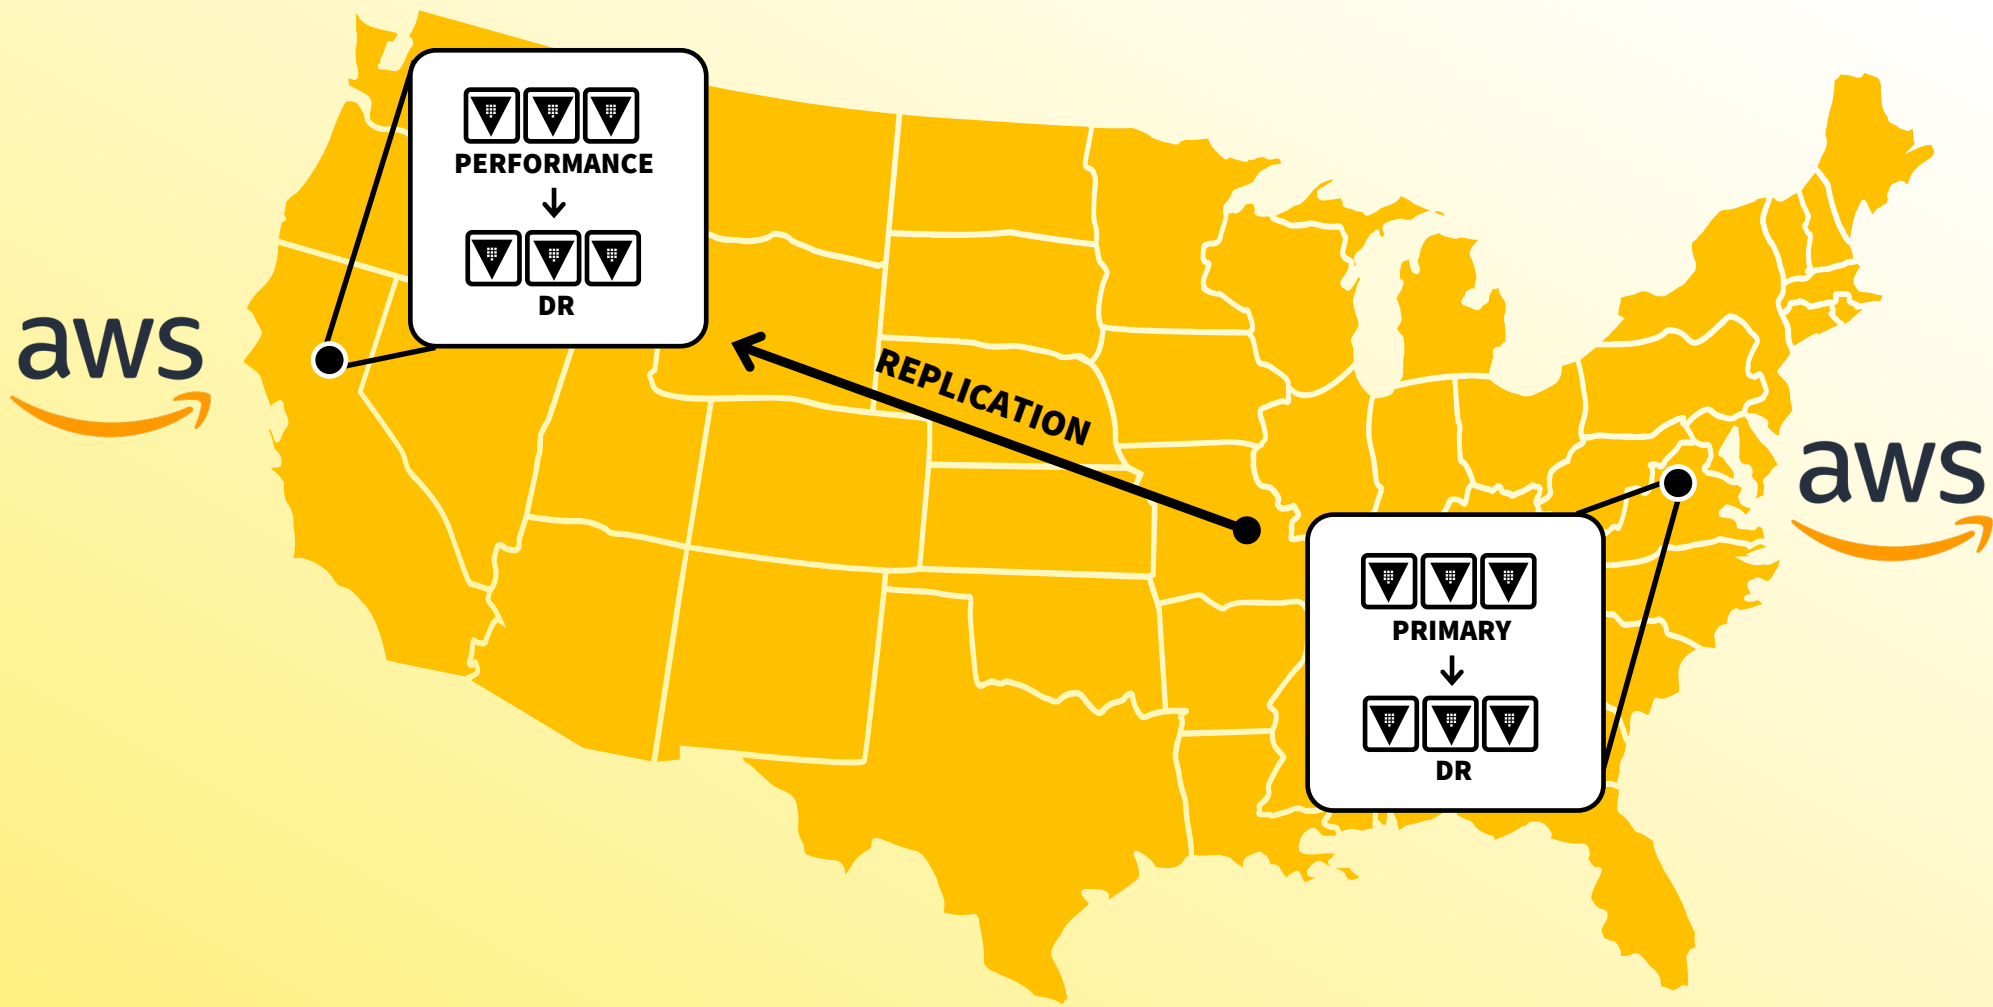
- Provides a warm-standby cluster where EVERYTHING is replicated to the DR secondary cluster(s)
- DR clusters DO NOT respond to clients unless they are promoted to a primary cluster
- Even as an admin or using a root token, most paths on a secondary cluster are disabled, meaning you can't do much of anything on a DR cluster

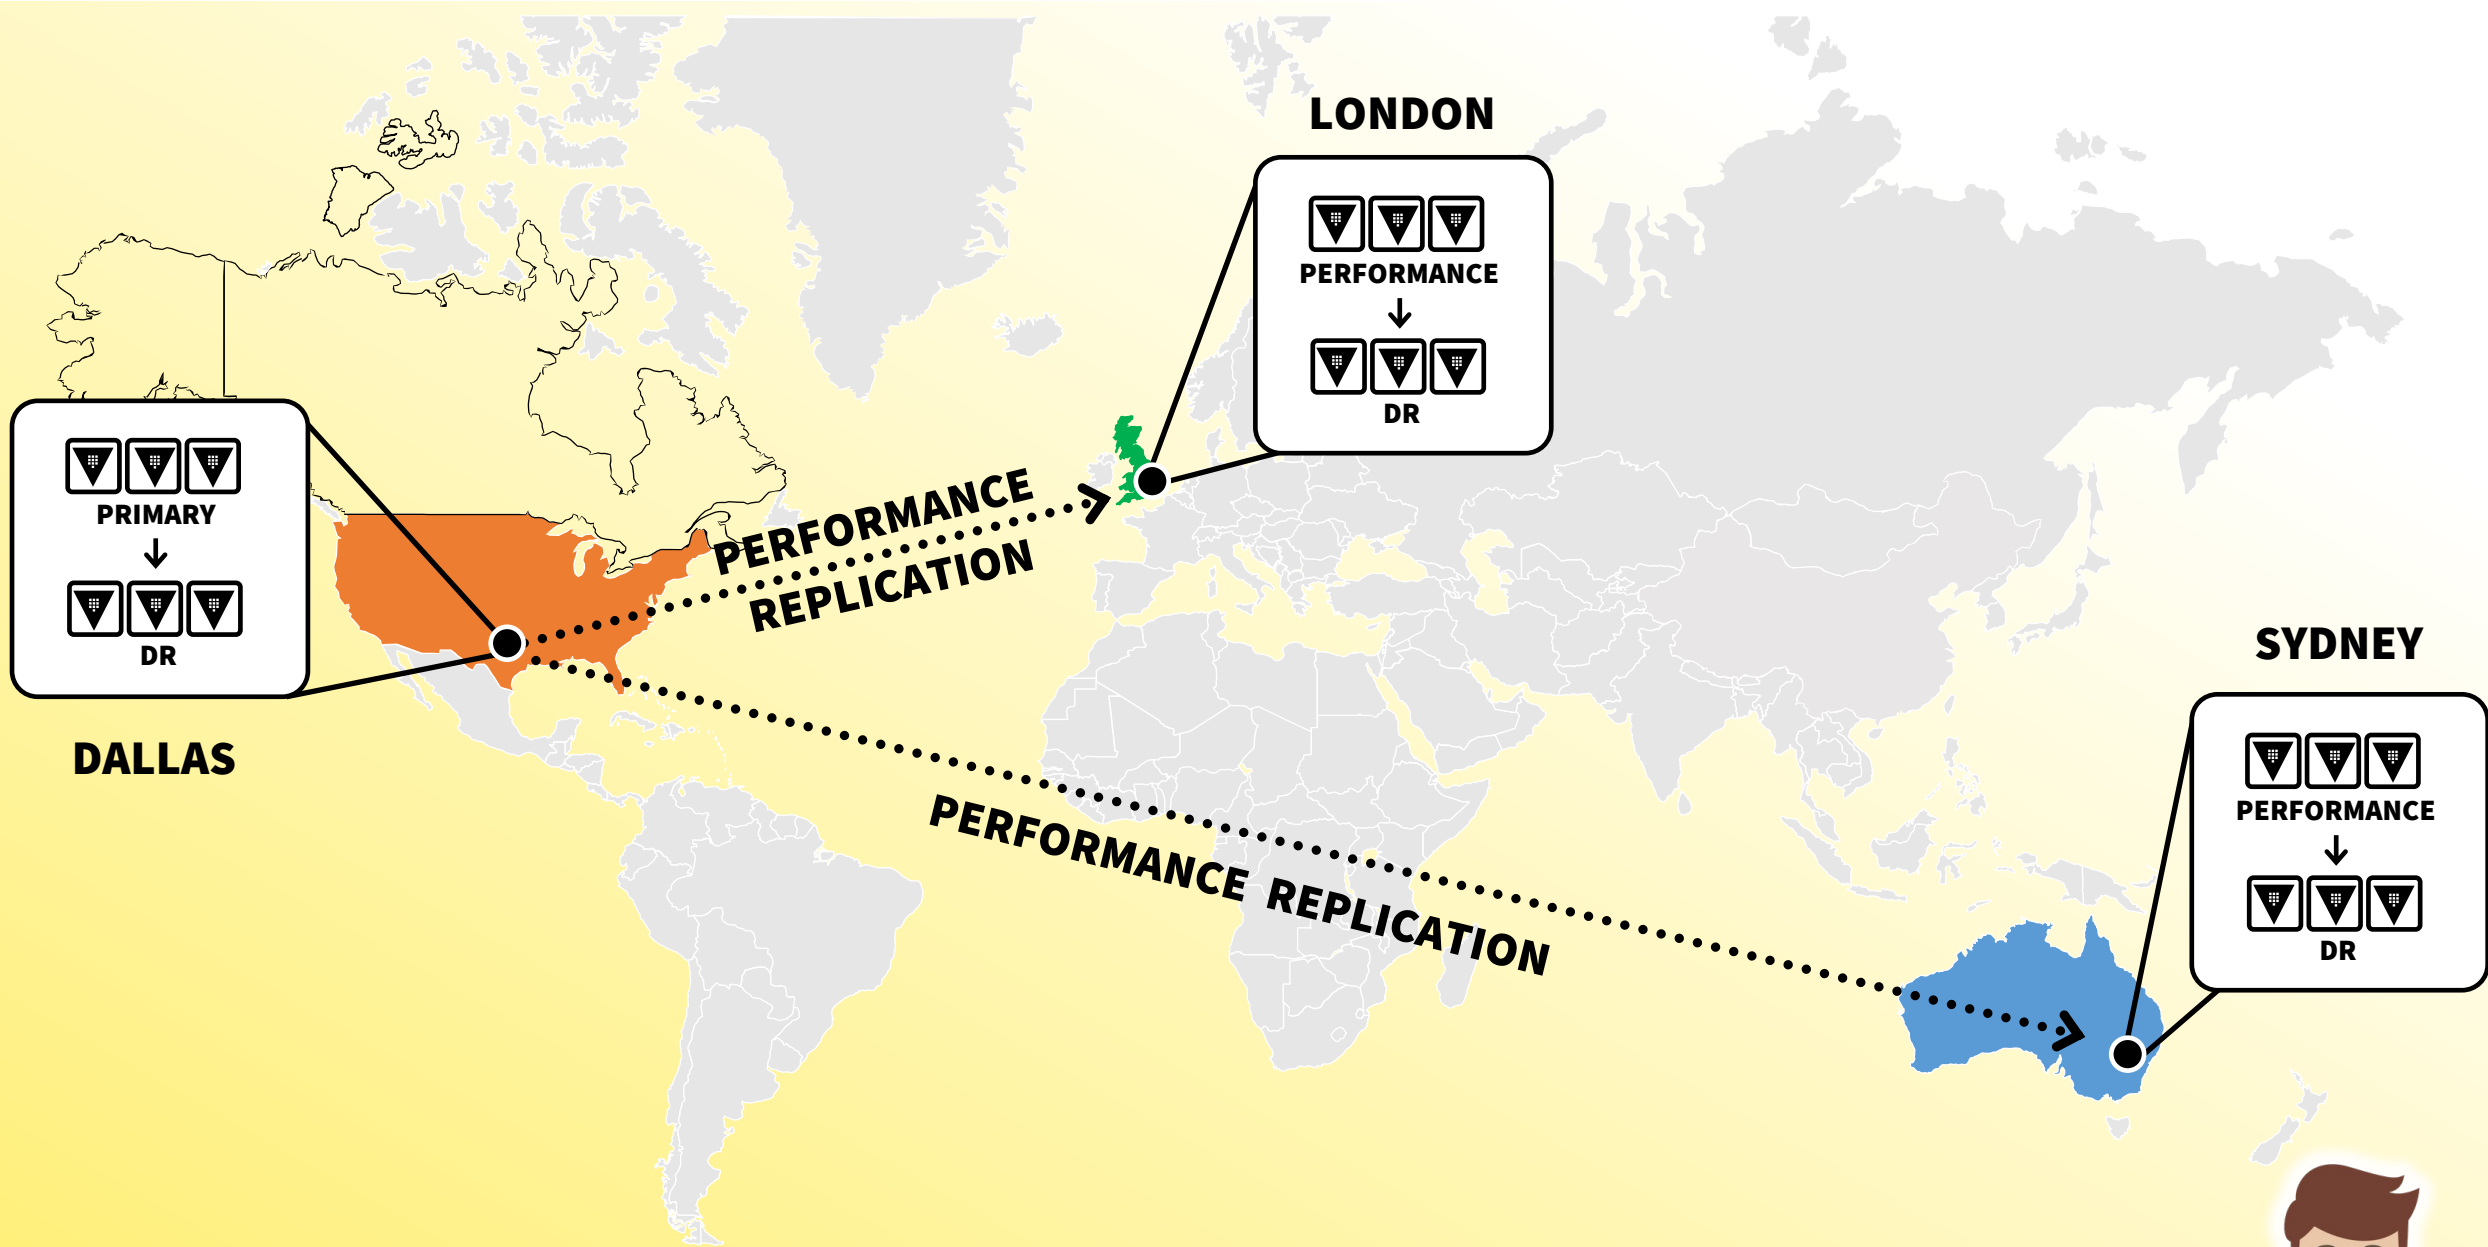


Replication Architecture




Replication Architecture

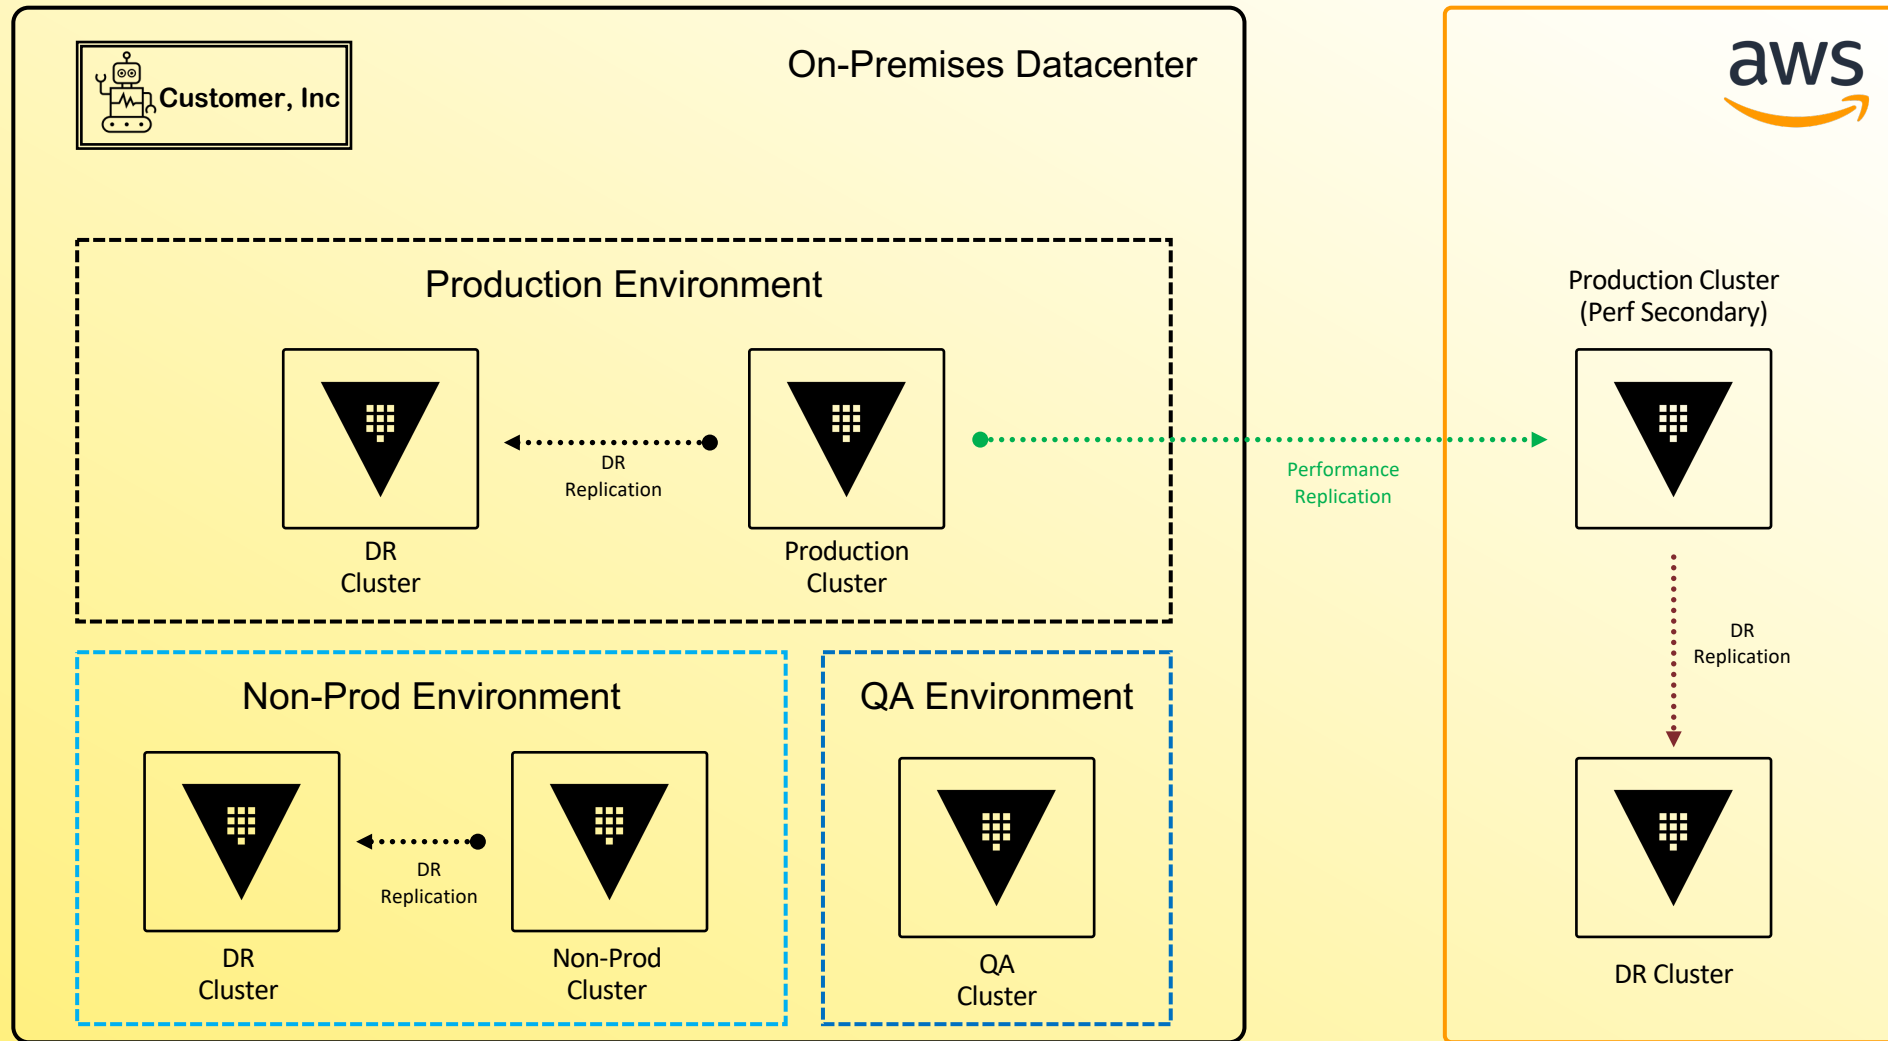




Real-World Customer Example




Real-World Customer Example



Networking Requirements



- Communication between clusters must be permitted to allow replication, RPC forwarding, and cluster bootstrapping to work as expected.
- If using DNS, each cluster must be able to resolve the name of the other cluster




Networking Ports



Source	Destination	Port	Protocol	Direction	Purpose
Client Machines	Load balancer	443	tcp	Incoming	Request distribution
Load Balancer	Vault Servers	8200	tcp	Incoming	Vault API
Vault Servers	Vault Servers	8200	tcp	Bidirectional	Cluster bootstrapping
Vault Servers	Vault Servers	8201	tcp	Bidirectional	Raft, replication, request forwarding
Vault Servers	External Systems	Various	Various	Various	External APIs



How Do We Set All of this Up?



Activate DR Replication
on the Primary as a DR
Primary

Create a secondary token
on the Primary cluster

Activate DR Replication
on the Secondary
cluster as a DR
secondary

Watch Vault replicated
the data from the
Primary to the new
Secondary cluster



Activating DR Replication



- Replication is NOT enabled by default, so you must enable it on each cluster that will participate in the replica set
- Enables an internal root CA on the primary Vault cluster - creates a root certificate and client cert
- Vault creates a mutual TLS connection between the nodes using self-signed certificates and keys from the internal CA – *NOT the same TLS configured for the listener*
 - If Vault sits behind a load balancer which is terminating TLS, it will break the mutual TLS between the nodes if inter-cluster traffic is forced through the load balancer



Secondary Token



- A secondary token is required to permit a secondary cluster to replicate from the primary cluster
- Due to its sensitivity, the secondary token is protected with response wrapping
- Multiple people should “have eyes” on the secondary token once it’s been issued until it is submitted to the secondary cluster
- Once the token is successfully used, it is useless (single-use token)
- The secondary token includes information such as:
 - The redirect address of the primary cluster
 - The client certificate and CA certificate



Secondary Token - Unwrapped

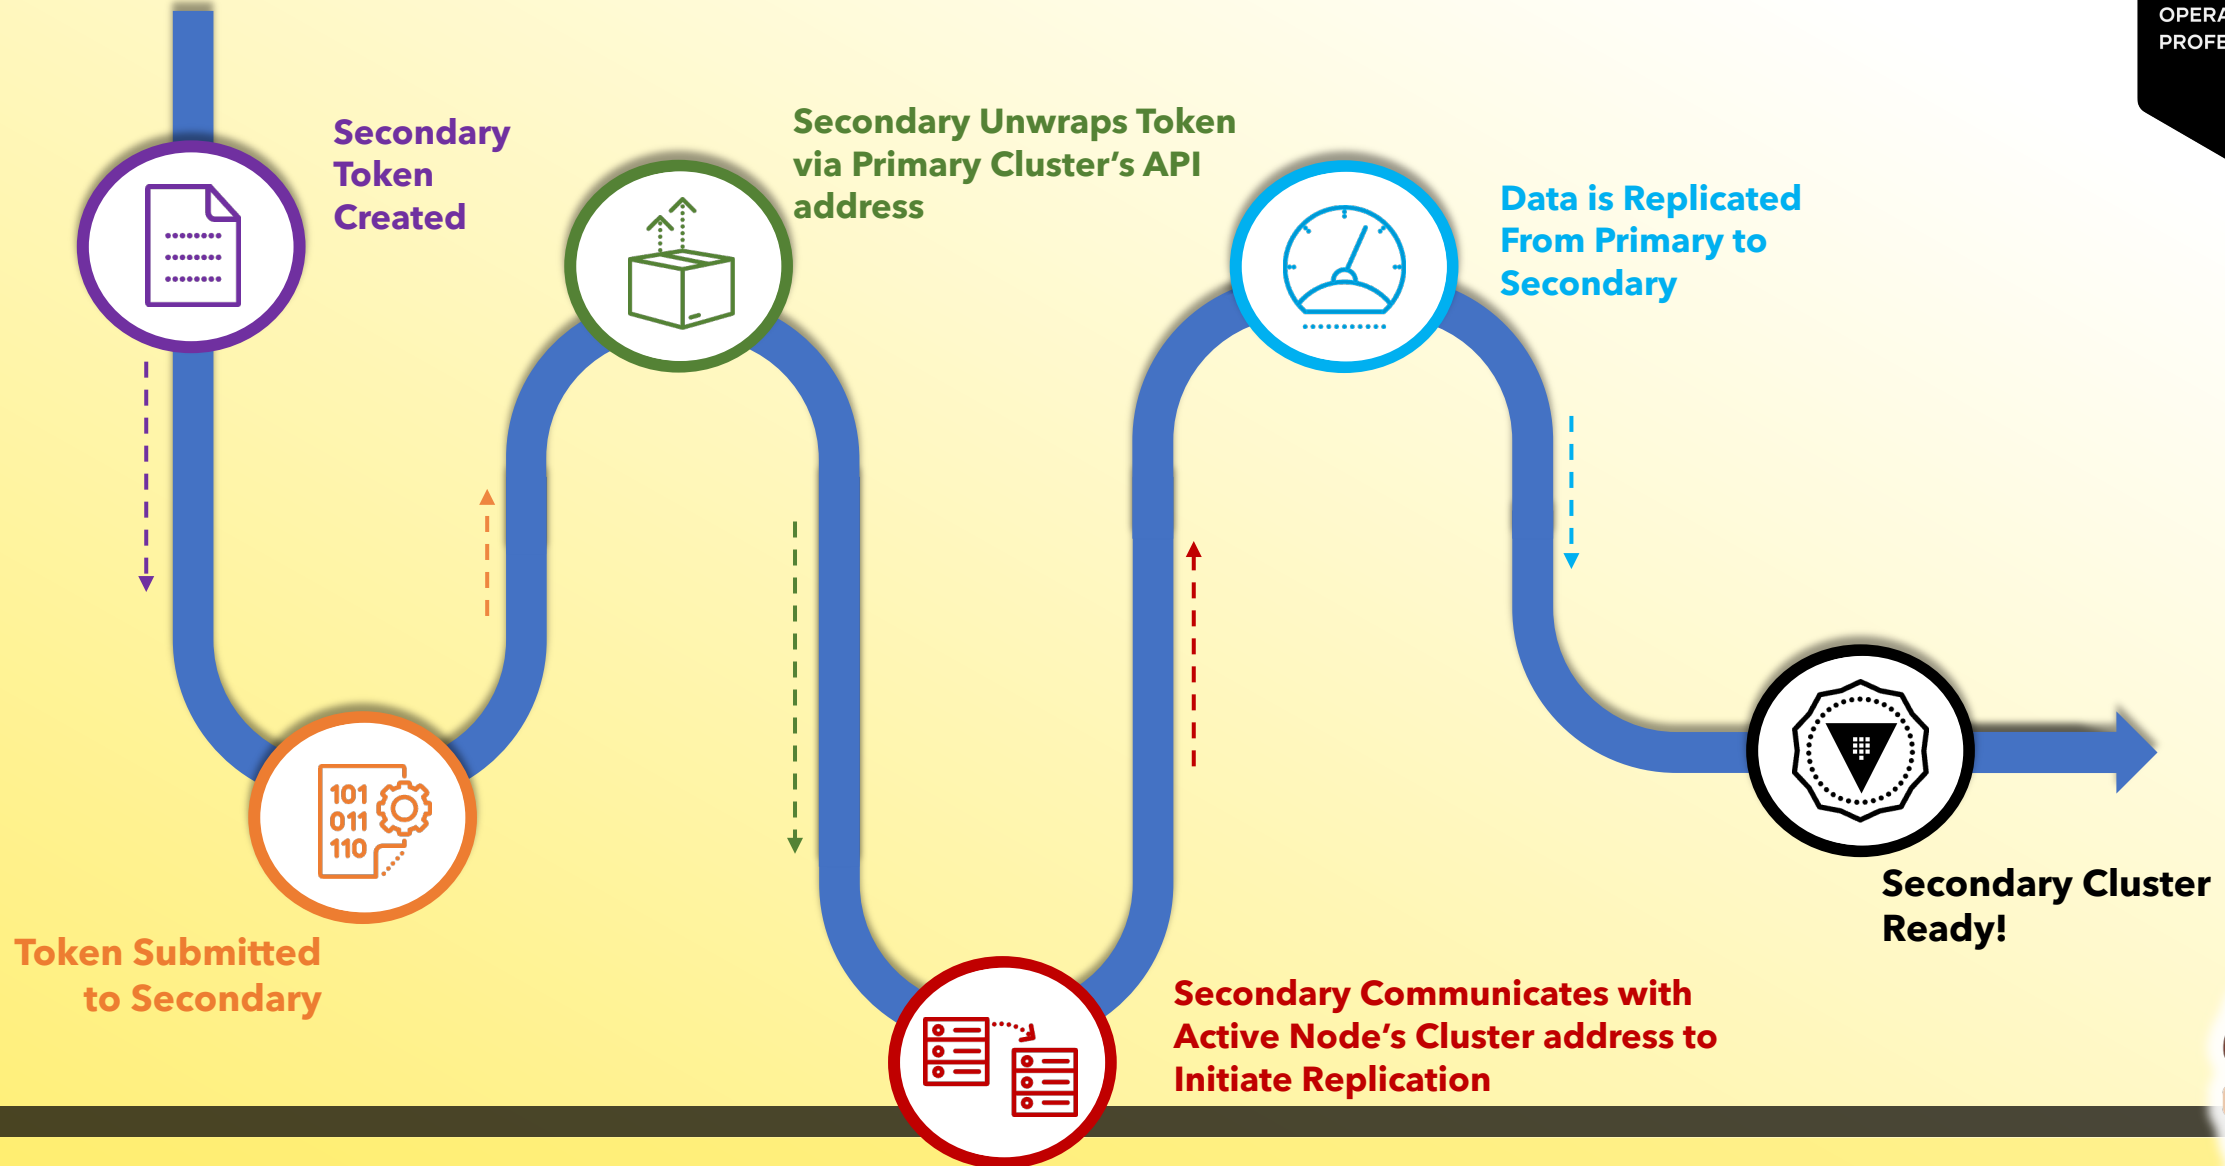


```
{
  "request_id": "98d4c7a5-0f00-4872-1cad-6ab8fa35694c",
  "lease_id": "",
  "lease_duration": 0,
  "renewable": false,
  "data": {
    "ca_cert":
    "MIICfjCCAd+gAwIBAgIIIVQciUM014jswCgYIKoZlZj0EAwQwMzExMC8GA1UEAxMocmVwLTA3MzQyYTBiLWJhZjktNTRhZC00MjcyLWVlZTE0NTFmMGQyNDAGFw0yMjA1MjMxNzIxMDUyMDUyMzA1MzE0OVowMzExMC8GA1UEAxMocmVwLTA3MzQyYTBiLWJhZjktNTRhZC00MjcyLWVlZTE0NTFmMGQyNDABmzAQBgqhkhjOPQIB...",
    "client_cert":
    "MIICZjCCAcigAwIBAgIIKW4DvMJIDt4wCgYIKoZlZj0EAwQwMzExMC8GA1UEAxMocmVwLTA3MzQyYTBiLWJhZjktNTRhZC00MjcyLWVlZTE0NTFmMGQyNDAGFw0yMjA1MjMxNzIxMDUyMDUyMzA1MzM1M1owLzEtMCSGA1UEAxMkZjYwNmEwMGltMTA0Ny05...",
    "client_key": {
      "d":
      1000631355517086513122196214347690053058610203119167515956358237211447177696212705845570913960352147412040118660857971566143956149412938809960381549100740826,
      "type": "p521",
      "x":
      6585241467240384151398124142600469244382875941120587428008118368573328804955608918211668669530795701495917170318651699823329298690163971349362335317686304875,
      "y":
      4563340717429320656179725289836652789047992587356159319649284729225610938283331963913484853756937351659805499727826936061640752374496368580488067455136501717
    },
    "cluster_id": "0d127970-99ce-152f-0311-3b081d1264d3",
    "encrypted_client_key": null,
    "id": "secondary",
    "mode": 512,
    "nonce": null,
    "primary_cluster_addr": "https://vault-pri.hcvop.com:8201",
    "primary_public_key": null
  },
  "warnings": null
}
```

This is not a normal thing you would do. I simply did it to show you what information the secondary token included



How is the Secondary Token Used?



Configure Replication on the CLI



1 Activate DR Replication

```
primary$ vault write -f sys/replication/dr/primary/enable
```

2 Create the Secondary Token

```
primary$ vault write sys/replication/dr/primary/secondary-token id=<id>
```

Name it what
you want

3 Activate the Secondary Cluster

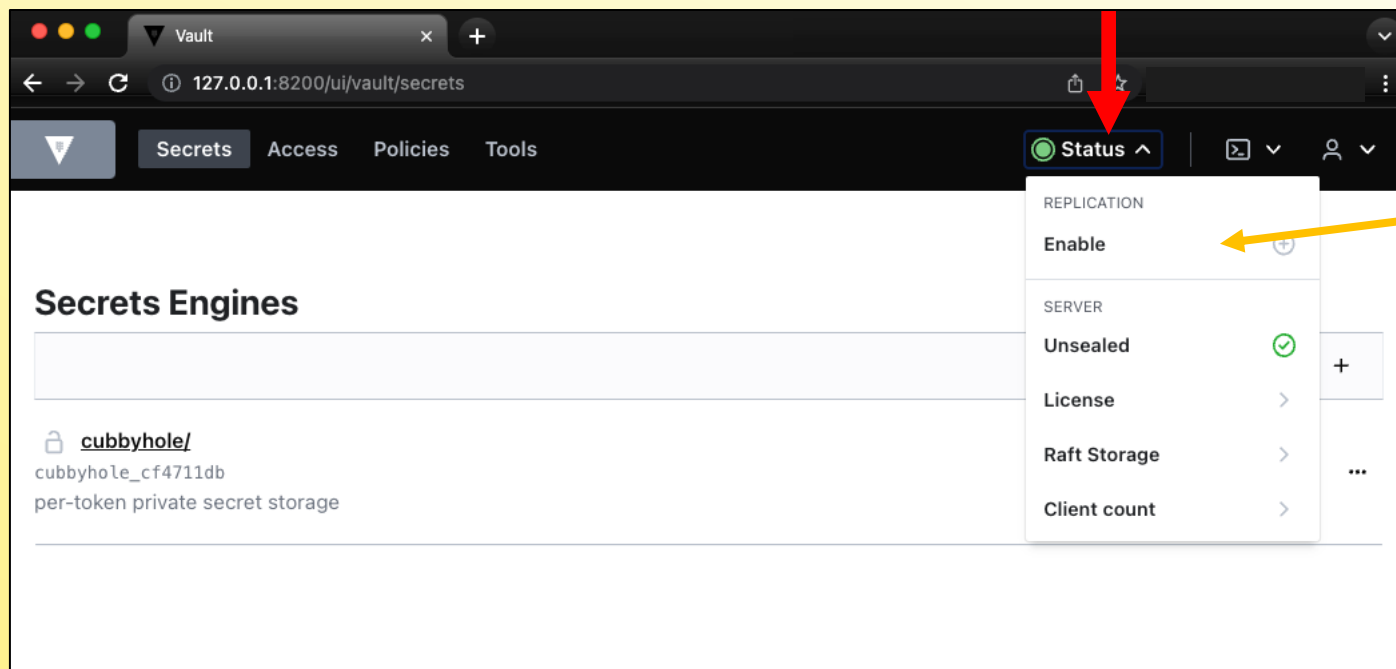
```
secondary$ vault write sys/replication/dr/secondary/enable token=<token>
```

Provide token from
primary cluster
(command above)



Configure Replication using the UI

Enable Replication on Primary

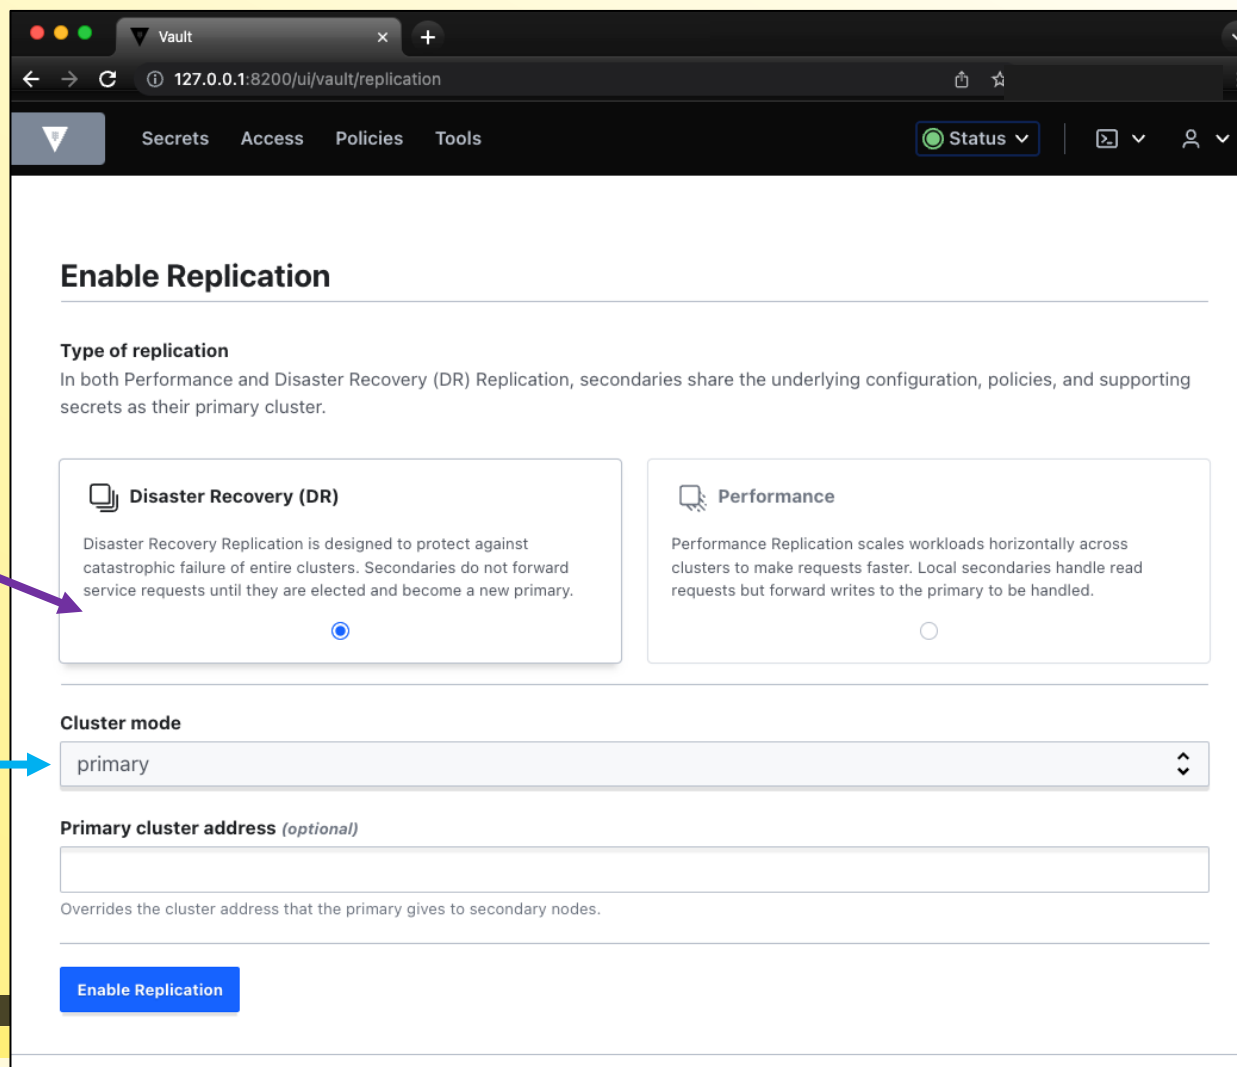


Enable
Replication



Configure Replication using the UI

Select Type and Mode on Primary



The screenshot shows the 'Enable Replication' page in the Vault UI. The browser address bar shows '127.0.0.1:8200/ui/vault/replication'. The page has a navigation bar with 'Secrets', 'Access', 'Policies', and 'Tools'. A 'Status' dropdown is visible. The main content area is titled 'Enable Replication'. Under 'Type of replication', there is a description: 'In both Performance and Disaster Recovery (DR) Replication, secondaries share the underlying configuration, policies, and supporting secrets as their primary cluster.' Below this are two options: 'Disaster Recovery (DR)' and 'Performance'. 'Disaster Recovery (DR)' is selected, indicated by a blue radio button. Below these options is a 'Cluster mode' dropdown menu set to 'primary'. There is also a 'Primary cluster address (optional)' text input field. At the bottom is a blue 'Enable Replication' button.

Enable Replication

Type of replication
In both Performance and Disaster Recovery (DR) Replication, secondaries share the underlying configuration, policies, and supporting secrets as their primary cluster.

Disaster Recovery (DR)
Disaster Recovery Replication is designed to protect against catastrophic failure of entire clusters. Secondaries do not forward service requests until they are elected and become a new primary.

Performance
Performance Replication scales workloads horizontally across clusters to make requests faster. Local secondaries handle read requests but forward writes to the primary to be handled.

Cluster mode
primary

Primary cluster address (optional)
Overrides the cluster address that the primary gives to secondary nodes.

Enable Replication

Select DR Replication

Choose Primary for the DR Primary cluster



Configure Replication using the UI

Add a Secondary



The screenshot shows the Vault UI for configuring replication. The browser address bar shows the URL `127.0.0.1:8200/ui/vault/replication/dr`. The navigation bar includes "Secrets", "Access", "Policies", and "Tools". The main content area is titled "Disaster Recovery" with a "primary" tag. Below the title are tabs for "Details", "Manage", and "Secondaries". A purple callout box labeled "View existing Secondaries" points to the "Secondaries" tab. On the left, there are two summary cards: "State" showing "running" with a green checkmark, and "Last WAL entry" showing the index "57". On the right, a section titled "0 Known Secondaries" contains the text "No known dr secondary clusters associated with this cluster" and a "View all" link. A red callout box labeled "Add a new Secondary" points to a red-bordered button labeled "Add secondary". At the bottom, a field for "primary_cluster_addr" is shown with the value "None set".



Configure Replication using the UI

Name Secondary and Get Secondary Token



The screenshot shows the Vault web interface in a browser window. The address bar shows the URL `127.0.0.1:8200/ui/vault/replication/dr/secondaries/add`. The navigation bar includes 'Secrets', 'Access', 'Policies', 'Tools', and a 'Status' dropdown. The main content area is titled 'Disaster Recovery' with a 'primary' tab. Below this are tabs for 'Details', 'Manage', and 'Secondaries'. The 'Secondaries' tab is active, showing a form to 'Generate a secondary'. The form includes a 'Secondary ID' input field with the value 'my-first-dr-secondary', a 'Time to Live (TTL) for generated secondary token' toggle (which is turned off), and 'Generate token' and 'Cancel' buttons at the bottom.

< Replication < Disaster Recovery

Disaster Recovery primary

Details Manage **Secondaries**

Generate a secondary

Generate a token to enable Disaster Recovery replication or change primaries on secondary cluster.

Secondary ID

This will be used to identify a secondary cluster once a connection has been established with the primary.

☐ **Time to Live (TTL) for generated secondary token**
If not set, the default value (30 minutes) will be used

Generate token **Cancel**

Give it a Name



Configure Replication using the UI

Copy New Secondary Token from Primary Cluster



Secondary
Token

Copy your token

This token can be used to enable Disaster Recovery replication or change primaries on the secondary cluster.

Activation token

```
eyJhbGciOiJIUzUxMiIsInR5cCI6IkpXVCJ9.eyJhY2Nlc3Nvcil6IiIsImFkZHIiOiJodHRwOi8vYnRrLW1hY2Jvb2stcHJvOjgyMDAiLCJleHAiOiJlE2NTMzMzA2MzgsImIhdCI6MTY1MzMzMyODgzOCwianRpljoiaHZzLkRYSWV5ZGw3aU5TVWRyMjIQUWZ1S0t6Rylsm5iZiil6MTY1MzMzMyODgzMywidHlwZSI6IndyYXBwaW50LnQALUNkeWVudAolOek82izumFEeafatE5NiIAJlv7506aueQ_0hPiAbQ
```

TTL 1800s

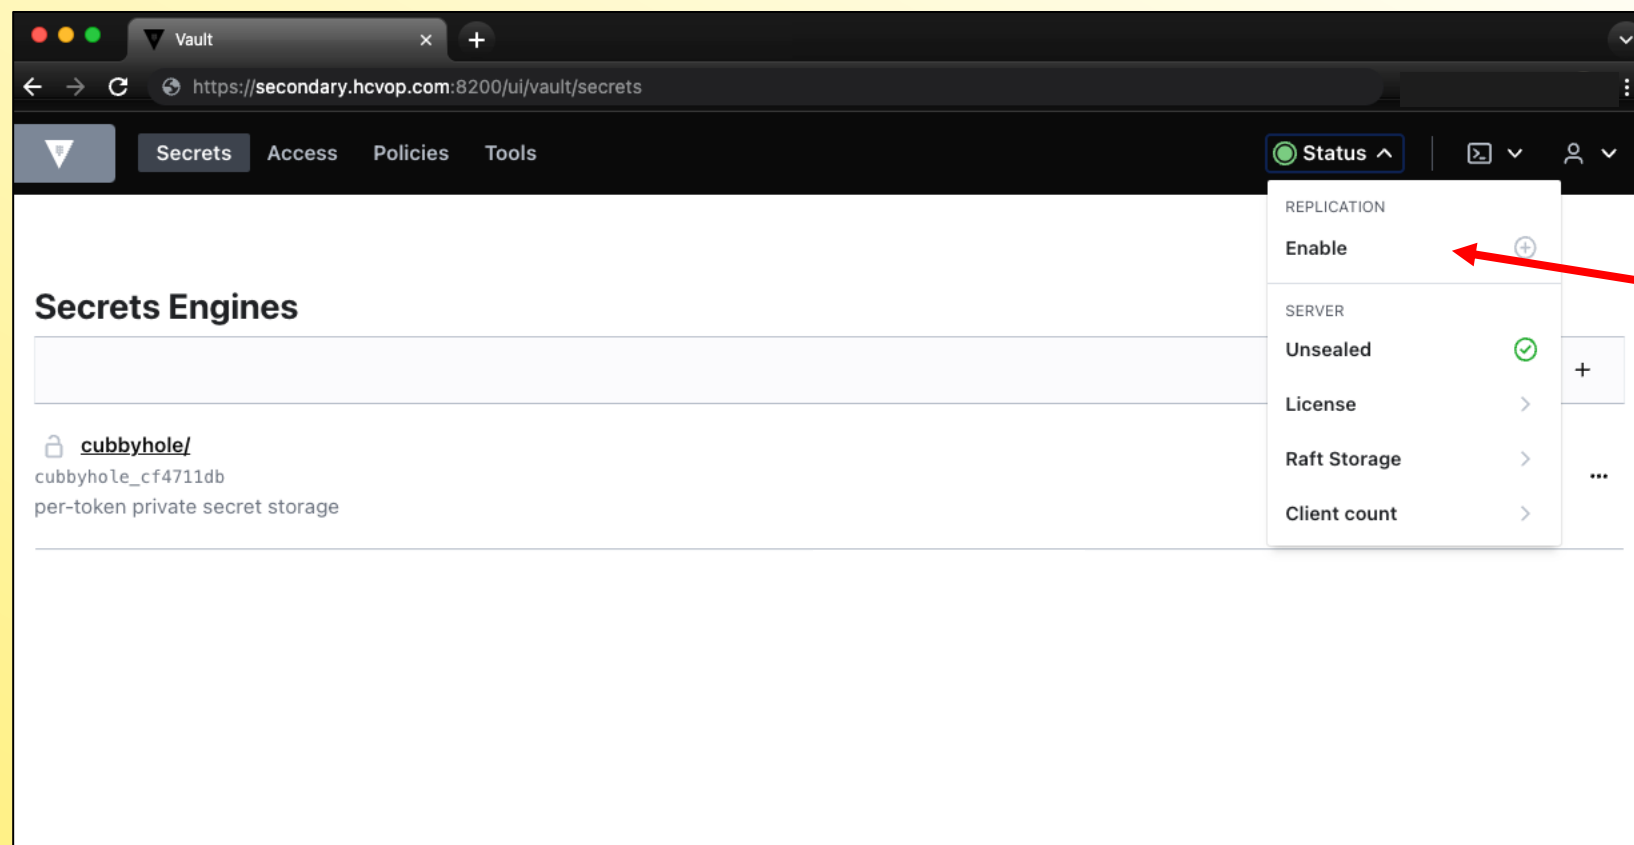
Expires

Copy & Close



Configure Replication using the UI

Enable Replication on Secondary Cluster

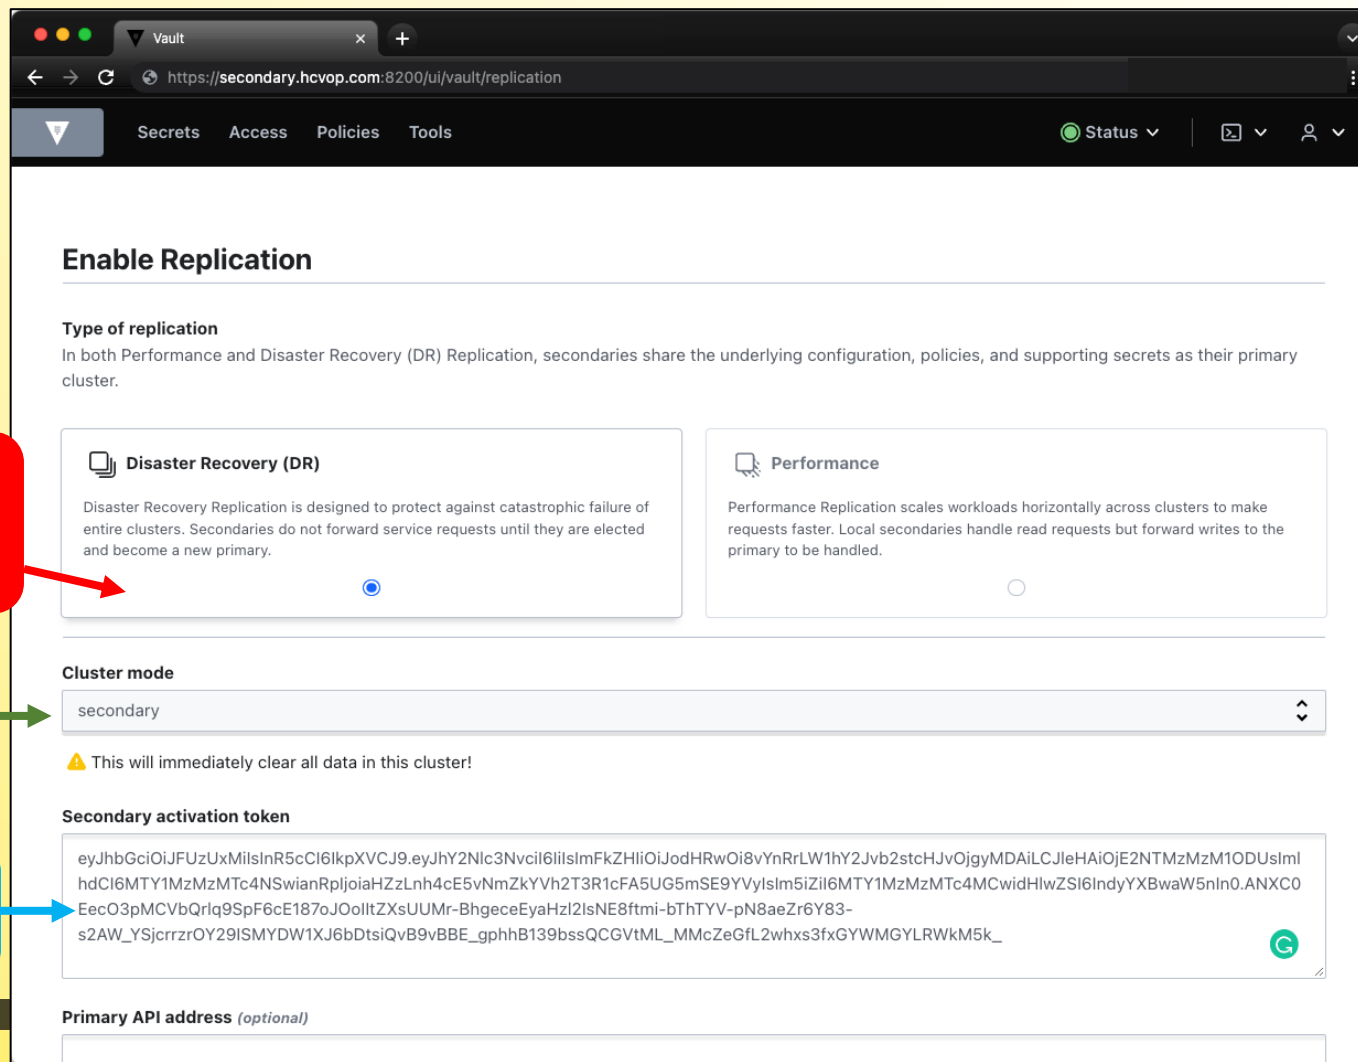


Enable
Replication on
Secondary



Configure Replication using the UI

Configure Secondary Cluster for Replication as a Secondary



The screenshot shows the Vault web interface for configuring replication. The browser address bar shows `https://secondary.hcvop.com:8200/ui/vault/replication`. The navigation bar includes "Secrets", "Access", "Policies", "Tools", and a "Status" dropdown. The main heading is "Enable Replication".

Type of replication
In both Performance and Disaster Recovery (DR) Replication, secondaries share the underlying configuration, policies, and supporting secrets as their primary cluster.

Two options are presented as cards:

- Disaster Recovery (DR)**: "Disaster Recovery Replication is designed to protect against catastrophic failure of entire clusters. Secondaries do not forward service requests until they are elected and become a new primary." This option is selected with a blue radio button.
- Performance**: "Performance Replication scales workloads horizontally across clusters to make requests faster. Local secondaries handle read requests but forward writes to the primary to be handled." This option is unselected with a grey radio button.

Cluster mode
A dropdown menu is set to "secondary". Below it is a warning icon and text: "This will immediately clear all data in this cluster!".

Secondary activation token
A text area contains a long alphanumeric token. A green copy icon is visible on the right side of the text area.

Primary API address (optional)
An empty text input field is located at the bottom of the form.

Select Disaster
Recovery
Replication

Select Secondary

Paste Secondary
Token Here



Monitor Replication



Check Status of ALL Replication

```
$ vault read -format=json sys/replication/status
```

Check Status of Performance Replication

```
$ vault read -format=json sys/replication/performance/status
```

Performance Replication Only

Check Status of DR Replication

```
$ vault read -format=json sys/replication/dr/status
```

DR Replication Only

